

ML2 – Probabilities; Bayesian decision theory

Stefano Rovetta

A.y. 2023-2024

Possible scenarios

1 (useful mostly for reasoning):

- Learner = hypothesis space = \mathcal{H} is fixed
- Data are fixed (population)

→ Find correct learners in \mathcal{H} (a *representation* task)

2 (not realizable):

- No data necessary: Probabilities are assumed to be known!

→ Find best hypothesis space \mathcal{H} and optimal learner

3 (your usual situation):

- Data will be **stochastic** but probabilities are **not known**
- Hypothesis space \mathcal{H} fixed, chosen in advance

→ Find learners in \mathcal{H} which are correct **for any possible realization of the data** (a *generalization* task)

Scenario 2:

Complete (probabilistic) knowledge

Probabilities

Probability

An expression of uncertainty.

- **Frequentist probability**

Probability of an event as a generalization of the frequency of occurrence of that event in infinite repetition of an experiment (**trial**).

- **Subjective probability**

Probability of an event as a confidence in the fact that the event itself will occur, even in a single experiment.

Probability

Something that may or may not happen is called an **outcome**

Call it ω

All possible outcomes constitute the **sample space** – the set $\Omega = \{\omega_1, \omega_2, \dots, \omega_N\}$ of all possible outcomes

E.g., results of an experiment, measures of a quantity, weather conditions...

We call **events** all possible subsets of Ω = combinations of possible outcomes:

Event space:

$$\begin{aligned} & \left\{ \emptyset, \{\omega_1\}, \{\omega_2\}, \dots, \{\omega_N\}, \{\omega_1, \omega_1\}, \{\omega_1, \omega_2\}, \dots, \{\omega_1, \omega_N\}, \{\omega_2, \omega_2\}, \{\omega_2, \omega_3\}, \dots, \Omega \right\} \\ &= \{A_1, A_2, A_3, \dots, A_{2^N}\} \end{aligned}$$

The **power set** of Ω , which has cardinality (size) 2^N

Examples

Tossing a coin

Outcomes: head H or tail $T \longrightarrow \Omega = \{H, T\}$

All events: $\{H\}, \{T\}, \{HT\}, \emptyset$

Only subsets containing one element are possible, the others will not occur:
Events here are **mutually exclusive**

Tossing a coin twice

$\Omega = \{HH, TT, TH, HT\}$

Some possible events:

- “tail only”: $A = \{TT\}$
- “at least one head”: $A = \{HT, TH, HH\}$
- “no two results are the same”: $A = \{HT, TH\}$

Some more examples

Weather forecasts

Outcomes: $\Omega = \{S \text{ sunny}, O \text{ overcast}, R \text{ rain} \}$
(mutually exclusive)


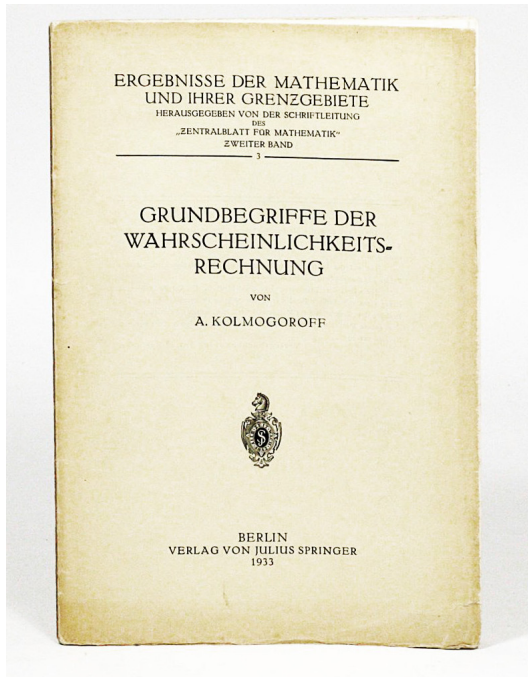
Weather forecasts for the next three days

Outcomes: $\Omega = \{SSS, SSO, SSR, SOS, SOO, SOR, \dots\}$ ($3^3 = 27$ combinations)

Possible event: today it is Wednesday; will it be sunny on Saturday?

$A = \{SSS, SOS, SRS, OSS, OOS, ORS, RSS, ROS, RRS\}$

Andrej N. Kolmogorov



Set-theoretic concepts apply

Summary of Terminology

Ω	sample space
ω	outcome (point or element)
A	event (subset of Ω)
A^c	complement of A (not A)
$A \cup B$	union (A or B)
$A \cap B$ or AB	intersection (A and B)
$A - B$	set difference (ω in A but not in B)
$A \subset B$	set inclusion
\emptyset	null event (always false)
Ω	true event (always true)

We assign a number to each event

$P(A)$ the probability of event A

Axioms of probability:

- 1 $P(A) \geq 0$
- 2 $\sum_{i=1}^N P(\omega_i) = 1$, or $P(\Omega) = 1$
- 3 If A_1 and A_2 are mutually exclusive events
(viewed as sets: if they are **disjoint** = have zero intersection),
then $P(A_1 \text{ or } A_2) = P(A_1 \cup A_2) = P(A_1) + P(A_2)$

More properties

(derived from the axioms)

- 1 $P(A) \leq 1$
- 2 If A_1 and A_2 are not mutually exclusive,
then $P(A_1 \cup A_2) = P(A_1) + P(A_2) - P(A_1 \cap A_2)$
- 3 $P(A_1 \cap A_2) = P(A_1)P(A_2)$ for *independent events*

two events are *independent* if the outcome of one event does not influence the outcome of the second event.

(OPPOSITE OF MUTUALLY EXCLUSIVE)

Example



$$\Omega^1 = \{1, 2, 3, 4, 5, 6\}$$

$$P(\omega_i) = P(\omega_j) \quad \forall i, j$$

$$\sum_i P(\omega_i) = P(\Omega^1) = 1 \Rightarrow P(\omega_i) = 1/6 \quad \forall i$$

Suppose we want to know how likely it is that the result is less than 3:

$$\text{Event} = \{1, 2\} \Rightarrow P(\omega < 3) = P(1 \cup 2) = P(1) + P(2) = 1/3$$

Example



$$\begin{aligned}\Omega^2 &= \Omega^1 \times \Omega^1 \\ &= \left\{ \{1, 1\}, \{1, 2\}, \{1, 3\}, \{1, 4\}, \{1, 5\}, \{1, 6\}, \{2, 1\}, \{2, 2\}, \{2, 3\}, \{2, 4\}, \dots \right\}\end{aligned}$$

Example (cont.)

We bet on 10. What is the probability of winning?

$$P(10) = P(\{4, 6\} \cup \{5, 5\} \cup \{6, 4\})$$


Outcomes are independent, so $P(\{\omega_i, \omega_j\}) = P(\{\omega_i\} \cap \{\omega_j\}) = P(\omega_i)P(\omega_j)$:

$$P(5 \cap 5) = P(5)P(5) = \frac{1}{6} \frac{1}{6} = \frac{1}{36}$$

Probability of any other pair $(6 \cap 4)$, $(4 \cap 6)$... is the same (fair dice)

$$\begin{aligned} P(10) &= (P(4)P(6)) + (P(5)P(5)) + (P(6)P(4)) = 3 \times \frac{1}{36} = \frac{1}{12} \\ &= \text{num. combinations} \times \left(\frac{1}{6}\right)^{\text{num. dice}} \end{aligned}$$

Note that...



...probabilities for different outcomes are not the same!

Continuous sample spaces

Examples: Age of a person, voltage in a circuit, force or torque in a mechanical link, temperature, time of the day...


All definitions still apply, at least in non-pathological situations

A trick for visualising probabilities:

- Think of probabilities for discrete sample and event spaces as **counts** and **frequencies**
- Think of probabilities for continuous sample and event spaces as **areas**, **volumes**, **hypervolumes**...

Further example


You draw one playing card from a full 52-card deck



- Outcomes: ω = one specific card
- Sample space Ω : The set of all 52 possible cards
- Event space 2^Ω : all possible types of cards (examples follow: compute probabilities for each!)


Some possible events

Bet on a specific value




Some possible events

Bet on a specific suit



Some possible events

Bet on a specific colour




Black
suits

Red
suits

Some possible events

Bet on whether the card is or is not a face card



Further possible events

The power set is of cardinality $2^{52} \approx 4.5 \times 10^{15} \dots$

Very many different bets are possible:

- Bet on any face card of a specific colour
- Bet on a card of odd/even value
- Bet on a card whose value is 2 or 7
- ...
- ... and – of course – bet on a specific card (individual outcomes are possible events)

Remark: Probabilities can be seen as quantities generated by functions of events (note that we already write them as such, $P(A)$!)

Random variables

A **random variable** is a numerical variable that doesn't have a fixed value, but changes according to a given probability law.

Example: random number generators in programming
(C/C++ `rand()` from `stdlib.h`/`cstdlib`, Matlab `rand`, Python `numpy.random.rand()`...)

Example: voltage measurements at the terminals of a heated resistor

Values may be real or discrete

Characterisation of probability of any type of random variable

Discrete or continuous

Definition

Let x indicate a random variable with values in \mathcal{X} .

Given a specific value $\hat{x} \in \mathcal{X}$,

$$P_{\mathcal{X}}(\hat{x}) = \Pr(x \leq \hat{x})$$

is the **cumulative distribution function** or simply **distribution function** of events in \mathcal{X} .

Characterisation of probability of discrete random variables

Definition

Let x indicate a random variable with values in a numerable set \mathcal{X} ,
e.g., an integer number in $\mathcal{X} = \{0, 1\}$.

Given a specific event $\hat{x} \in \mathcal{X}$,

$$F_{\mathcal{X}}(\hat{x}) = \Pr(x = \hat{x})$$

is the **probability mass function** of \mathcal{X} .

Gives the finite probability that a random event x has a specific value \hat{x} .

Characterisation of probability of continuous random variables

Definition

Let x indicate a random variable with values in a non-numerable set \mathcal{X} ,
e.g., a real number in $\mathcal{X} = [0, 1]$.

A function $f_{\mathcal{X}}$ such that

$$P_{\mathcal{X}}(\hat{x}) = \int_{-\infty}^{\hat{x}} f_{\mathcal{X}}(x) dx$$

is the **probability density function** of \mathcal{X} .

Gives the **infinitesimal** probability that a random event x has value \hat{x} . (Infinitesimal = null)

Its **definite integral on a random interval** $[\hat{t}_1, \hat{t}_2] \in \mathcal{E}$ is the **finite** probability that $\hat{t}_1 < e < \hat{t}_2$.

Expectation

An important use of probability functions:
compute the “most likely” or **expected** value of some random X .
In the case of discrete events, it is a weighted sum

$$E\{X\} = \sum_i \xi_i F_X(\xi_i)$$

ξ_i are the possible values of X

The symbol $E\{\}$ is the expectation operator

Expectation

For real-valued X :

$$\mathbb{E}\{X\} = \int_{\mathcal{X}} \xi f_x(\xi) d\xi$$

Conditional probability

$P(E|F)$ — The probability of an event E **given** the knowledge that another event F has occurred

Example (dice): $P('10') = 1/12$

But if we know that the first dice is '2', then $P('10' \mid \text{first dice is '2'}) = 0$

If we know that the first dice is '5', then $P('10' \mid \text{first dice is '5'}) = 1/6$

Bayesian probabilities

Bayesian probability jargon

t is a hypothesis (event)

t_i is a set of alternative hypotheses, $t \in \{t_i\}$

x is an experimental observation

$P(t)$ is the **a priori probability** of hypothesis t :
probability that t is true before seeing any experimental observation


$P(t|x)$ is the **a posteriori probability** of hypothesis H after observing X

$P(x|t)$ is the **likelihood** of observing x when t holds (is verified, is true, is certain)

$P(x)$ is the **marginal probability** of x , the probability of observing x in any case

Total probability theorem

How to compute a marginal probability:



$$P(x) = \sum_i P(x|t_i)P(t_i)$$

Bayes' theorem

$$P(t_i|x) = \frac{P(x|t_i)P(t_i)}{P(x)}$$

Gives the probability of hypothesis t_i after seeing an experimental observation x

Bayes' theorem, alternate form

$$P(t_i|x) = \frac{P(x|t_i)P(t_i)}{\sum_{j=1}^c P(x|t_j)P(t_j)}$$

- Uses the total probability theorem
- Does not require $P(x)$
- The denominator is also known as a **partition function**.
It acts as a **normalizer** that makes the sum of all $P(t_i|x)$ (for $i : 1 \dots c$) equal one.

Bayesian Decision Theory

Pattern Classification, 2nd Edition

Richard O. Duda, Peter E. Hart, David G. Stork

ISBN: 978-0-471-05669-0

Wiley-Interscience

680 pages

November 2000

The decision problem

- c possible, mutually exclusive events, or “states of nature”

$$\{t_1, t_2, \dots, t_c\}$$

- s possible actions or “decisions” that we may make

$$\{y_1, y_2, \dots, y_s\}$$

We want to find a rule that, given any state t , makes the most suitable decision y .
The problem is that t may not be observable!

The decision process

From simplest to most complex:

- 1 Decisions:
states of nature \longrightarrow **actions**
- 2 Uncertainty in the states of nature:
probabilities \longrightarrow **states of nature** \longrightarrow **actions**
- 3 Conditional decisions:
observations \longrightarrow **probabilities** \longrightarrow **states of nature** \longrightarrow **actions**
- 4 Using Bayes formula:
observations \rightarrow [BAYES] \rightarrow **probabilities** \longrightarrow **states of nature** \longrightarrow **actions**

1: Decisions

states of nature \longrightarrow **actions**

Rationale

- Given the **state of nature** t , we act consequently and make a decision, or take an action, y .
- So the decision or action is a function of the state of nature, $y(t)$.
- Every decision has a **cost** – a high-cost decision is a “wrong” decision
- We measure this cost using a **loss function** $\lambda(y, t)$
- The loss function evaluates the cost of **each decision**, depending on the **true state of nature**, including our **subjective** considerations (preferences) as well as **objective** elements (actual costs)
- We build a **decision rule**, the function $y(t)$ that given t produces y , attempting to **minimise the loss**.

Example

Buying a pair of shoes. I have two choices:

- (Italian) size 43, cost 200
- Size 42, cost 120

Depending on my shoe size, a possible loss function is:

$$\lambda(y, t) = \lambda \left(\boxed{\text{what I buy}}, \boxed{\text{my actual size}} \right) =$$

	Buy 43	Buy 42
I have size 43	200	120 + uncomfortable
I have size 42	200 + uncomfortable	120

Decision rule that minimises the loss:

- If you have size 42, buy the cheaper pair which is 42
- If you have size 43... depending how you quantify “uncomfortable”, you
 - buy 42 (if you value “uncomfortable” less than 80)
 - or 43 (otherwise).

2: Uncertainty in the states of nature

probabilities \longrightarrow **states of nature** \longrightarrow **actions**

Probabilistic modelling

Usually t is not known with certainty at the time of making a decision.

Example: Insurance

The insurance company must establish the reimbursement policies *before* accidents happen.

Example: Measuring instrument

Due to measurement errors, the reading of an instrument *is not* necessarily the true value of a physical quantity.

⇒ We use probabilities.

$$\begin{array}{cccccc}
 t_1 & t_2 & t_3 & \dots & t_c \\
 P(t_1) & P(t_2) & P(t_3) & \dots & P(t_c)
 \end{array}$$

The cost of a decision now cannot be known with certainty

⇒ To evaluate each possible decision, we use the **expected** loss
(average over all possibilities, weighted with the respective probabilities)

For decision y_1 :

$$\begin{aligned} R(y_1) &= \lambda(y_1, t_1)P(t_1) + \lambda(y_1, t_2)P(t_2) + \dots + \lambda(y_1, t_c)P(t_c) \\ &= \sum_{j=1}^c \lambda(y_1, t_j)P(t_j) \end{aligned}$$

For a generic decision y_i :

$$R(y_i) = \sum_{j=1}^c \lambda(y_i, t_j)P(t_j)$$

$R(y_i)$ is the **risk** (expected loss) of decision y_i

The decision rule

In the presence of uncertainty, the decision rule is as before

But this time it must minimise the **risk** of the decision

In other words, it must minimise the loss **on average**

3: Conditional decisions

observations \longrightarrow probabilities \longrightarrow states of nature \longrightarrow actions

Unobservable quantities

Many interesting quantities cannot be measured directly.

Example: Disease

A doctor can measure sign and symptoms, but not directly the disease. Given (for instance) fever, this indicates that there *may be* a certain disease (e.g. flu). But flu itself cannot be measured with an instrument.

Example: Stock market trend

There are several indicators that the stock market *may go* in a certain direction in the future (increasing or decreasing), but the true dynamics of the stock market cannot be modelled.

⇒ We use **experimental observations** and **conditional probabilities**.

X

t_1

t_2

t_3

\dots

t_c

$P(t_1|\mathbf{x})$ $P(t_2|\mathbf{x})$ $P(t_3|\mathbf{x})$ \dots $P(t_c|\mathbf{x})$

Note: observations may be

- a scalar x
(for instance an individual sensor reading)
- a vector \mathbf{x}
(for instance all the pixels in an image)


Conditional risk

$$R(y_i | \mathbf{x}) = \sum_{j=1}^c \lambda(y_i, t_j) P(t_j | \mathbf{x})$$

$R(y_i | \mathbf{x})$ is the **conditional risk** of decision y_i
when we have the **experimental observation** \mathbf{x}

General structure of a decision rule

In the most general case we are considering, we have to design:



A system that, given any observation, outputs the best decision

The decision rule

The decision rule is more or less as before.

But this time, it should minimise the conditional risk **for all possible observations x !**

Of course this may not be possible in all cases.

So the realistic criterion is:

The decision rule must minimise the **average (expected) risk**
over all possible observations

So we must take the expectation of the risk over the observations as our criterion to be minimised.

Expected risk

If observations are discrete:

$$R = \sum_{\mathbf{x} \in \mathcal{X}} R(y(\mathbf{x}) | \mathbf{x}) P(\mathbf{x})$$

where $P(\mathbf{x})$ is the probability mass function of experimental observations and \mathcal{X} is the set of all possible inputs (the “input space”).

If observations are continuous:

$$R = \int_{\mathcal{X}} R(y(\mathbf{x}) | \mathbf{x}) p(\mathbf{x}) d\mathbf{x}$$

where $p(\mathbf{x})$ is the probability density function of experimental observations.

4: Using Bayes formula

observations \rightarrow [BAYES] \rightarrow **probabilities** \longrightarrow **states of nature** \longrightarrow **actions**

Bayes decision theory




Example: A doctor should diagnose a disease after visiting a patient

He records all observation and measurements into a patient record \mathbf{x}

Bayes decision theory

Usually a doctor has some information available from his medicine textbooks and from his own experience:

- The incidence of diseases
 $\rightarrow P(t_i)$
- The typical and not-so-typical signs and symptoms of diseases
 $\rightarrow P(\mathbf{x}|t_i)$




Bayes decision theory

From Bayes' theorem:

$$P(t_i|\mathbf{x}) = \frac{P(\mathbf{x}|t_i)P(t_i)}{\sum_{j=1}^c P(\mathbf{x}|t_j)P(t_j)}$$

Numeric example

Autonomous rover for unmanned explorations, equipped with four sensors



Numeric example

cont.

- **Possible states of nature:**

$$t \in \{t_1, t_2, t_3\} = \{\text{'water'}, \text{'solid ground'}, \text{'sand'}\}$$

- **Their a priori probabilities:**

$$P(t_1) = .2, P(t_2) = .4, P(t_3) = .4$$

- **Possible decisions:**

$$y \in \{y_1, y_2\} = \{\text{rover:retract}, \text{rover:advance}\}$$

- **Input observations (readings from sensors):**

$$\mathbf{x} = [\text{groundTouchSens1}, \text{groundTouchSens2}, \text{groundOptSens}, \text{groundWhisker}]$$

- **We now receive an input observation \mathbf{x} for which likelihoods are:**

$$P(\mathbf{x}|t_1) = .5, P(\mathbf{x}|t_2) = .9, P(\mathbf{x}|t_3) = .1$$

- **NOTE that LIKELIHOODS MAY NOT SUM UP TO 1**

They are not mutually exclusive

(“in direct competition with each other”)

Example

cont.

- We are given this loss matrix:

$$\Lambda = \begin{bmatrix} 0.1 & 1.0 & 4.0 \\ 2.0 & 0.1 & 0.1 \end{bmatrix}$$

- Conditional risk of decision y_1 given observation \mathbf{x} :

$$R(y_1|\mathbf{x}) = \sum_{j=1}^3 \lambda_{1j} P(t_j|\mathbf{x}) = 0.1 \times 0.5 + 1 \times 0.9 + 4 \times 0.5 = 2.95$$

- Conditional risk of decision y_2 given observation \mathbf{x} :

$$R(y_2|\mathbf{x}) = \sum_{j=1}^3 \lambda_{2j} P(t_j|\mathbf{x}) = 2 \times 0.5 + 0.1 \times 0.9 + 0.1 \times 0.5 = 1.14$$

Example

cont.


When we receive input \mathbf{x} ,
the decision that minimizes the conditional risk is y_2

The Bayes decision criterion: in short

*To minimize R , given an input \mathbf{x} the decision rule $y(\mathbf{x})$ should output the decision y that minimizes the **expected risk** R .*

Theoretically optimal criterion (you can't do better than this)

Errors are always possible!



1 and 2 are two example posterior probabilities

The Bayes error, the best possible error probability.

Classification

Classification is a decision problem with:

- $s \equiv c$
- $\{y_1, \dots, y_s\} \equiv \{t_1, \dots, t_c\}$

I.e. there is no actual decision to take, we are only recognizing the state of nature
(the **class**)

A loss function for classification

zero-one loss:

$$\lambda(y|t) = \begin{cases} 0 & \text{if } y = t \\ 1 & \text{if } y \neq t \end{cases}$$

Example: zero-one loss matrix for a three class problem ($c = 3$):

$$\Lambda = \begin{bmatrix} 0 & 1 & 1 \\ 1 & 0 & 1 \\ 1 & 1 & 0 \end{bmatrix}$$

What does the 0-1 loss mean?

- All types of errors have the same cost ($= 1$)
- Correct classifications don't have a cost
- $\Rightarrow R$ equals expected probability of error
(proof: plug zeroes and ones in the definition of conditional cost)

Zero-one loss = minimum-error-rate classification

Designing a classifier

A classifier is a rule $y()$ that receives an observation \mathbf{x} and outputs a **class** $y(\mathbf{x})$.


The Bayes decision criterion states that y should minimize $R(y(\mathbf{x}) | \mathbf{x})$

A natural idea:

- Build c blocks or “matched filters” $g_j, j : 1 \dots c$ that compute $g_1(\mathbf{x}) = -R(y = t_1 | \mathbf{x}), \dots, g_c(\mathbf{x}) = -R(y = t_c | \mathbf{x})$
- Select $y = t_j$ that has maximum $g_j(\mathbf{x})$

$g_j()$ are called **discriminant functions**

The operation of looking for the location j of the maximum value (the “argument” of the maximum) is called **argmax**



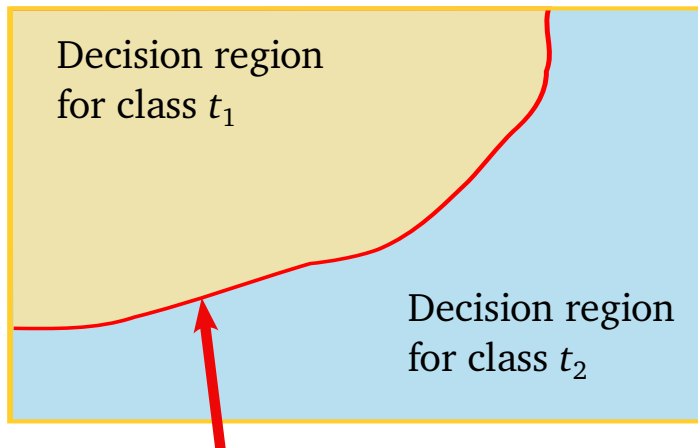
Decision regions

A **decision region** is a subset of the data space with a given minimum-conditional-risk decision (i.e., the decision y is the same for all data in the region)

Decision regions are separated by **decision boundaries** (or **decision surfaces**)

The decision boundary between two regions (say $y = t_j$ and $y = t_k$) are defined by:

$$g_j(\mathbf{x}) = g_k(\mathbf{x})$$



Decision boundary
between class t_2
and class t_2

Discriminant functions for zero-one loss

In the case of the minimum-error-rate classifier (= using zero-one loss):

$$\begin{aligned} g_i(\mathbf{x}) &= - \sum_{j=1, j \neq i}^c P(t_j|\mathbf{x}) \\ &= P(t_i|\mathbf{x}) - 1 \end{aligned}$$

where $P(t_i|\mathbf{x})$ is obtained from Bayes' theorem

Other ways to define discriminant functions

A classifier is defined by the decision boundaries
so the actual functions being compared need not be actually $g_j(\mathbf{x}) = -R(t_j|\mathbf{x})$

They can be any **monotonically increasing transformation** $g_j(\mathbf{x}) = f(-R(t_j|\mathbf{x}))$
that preserves decision boundaries

e.g., $g_j = \log R(t_j|\mathbf{x})$ or $g_j = \frac{1}{1+e^{-R(t_j|\mathbf{x})}}$

This gives us more freedom in building a classifier!

We can use more general **scores** $f(-R(t_j|\mathbf{x}))$
instead of actual conditional risks $-R(t_j|\mathbf{x})$.

Reasonable discriminant functions for zero-one loss

The transformation $f(x) = x + 1$ is monotonic and preserves decision boundaries, so we can avoid the useless -1 :

$$\begin{aligned} g_i(\mathbf{x}) &= 1 - \sum_{j=1, j \neq i}^c P(t_j|\mathbf{x}) \\ &= P(t_i|\mathbf{x}) \end{aligned}$$

Here we see that if we give the same weight to all errors (0/1 loss) the discriminant functions are simply the probability of each class given the input — so we take the one with maximum probability!

Quite sensible criterion.

A popular classifier: Naive Bayes

Is built using “wrong” discriminant functions based on “naive,” or even “idiot,” assumptions (→ also “Idiot’s Bayes Classifier”).

- **Recall that:** $\mathbf{x} = [x_1, x_2, x_3, \dots, x_d]$
- Let’s focus on discrete x s
- **So:** $P(t_i|\mathbf{x}) \propto P(\mathbf{x}|t_i)P(t_i) = P(x_1, x_2, x_3, \dots, x_d | t_i)P(t_i)$
- **In general:** $\Pr(a, b|c) \neq \Pr(a|c)\Pr(b|c)$
- **Naive assumption:** $\Pr(x_1, \dots, x_d | t_i) = \Pr(x_1 | t_i)\Pr(x_2 | t_i) \cdots \Pr(x_d | t_i)$

We are pretending that input variables are **all independent of each other**

Naive Bayes classifier

$$\begin{aligned} g_i(\mathbf{x}) &= P(t_i) [P(x_1|t_i) \times P(x_2|t_i) \times \dots \times P(x_d|t_i)] \\ &= P(t_i) \prod_{j=1}^d P(x_j|t_i) \end{aligned}$$

How a naive Bayes classifier “learns”

Particularly handy when features are binary (true/false):

To “learn” $P(x_k|t_i)$ we count how often each value of x_k occurs in class t_i in the training set:

$$P(x_k = \text{true}|t_i) = \frac{\text{number of times } x_k = \text{true in class } t_i}{\text{number of observations of class } t_i} = \frac{N_{\text{true},t_i}}{N_{t_i}}$$

$$P(x_k = \text{false}|t_i) = 1 - P(x_k = \text{true}|t_i) = 1 - \frac{N_{\text{true},t_i}}{N_{t_i}}$$

Prior probability of classes:

$$P(t_i) = \frac{\text{number of observations in class } t_i}{\text{number of observations in the training set}} = \frac{N_{t_i}}{N}$$

Not so idiot?

With many features the naive assumption is approximately correct!

Example: Spam detection

- Observations: email messages
- Features:
 - Presence of words typical of spam (from a list)
 - Presence of specific spelling mistakes
 - Mismatch between address shown in links and address actually pointed to
 - Only images and no text
 - Only attachments and little text
 - ...
- Training set: Your “JUNK” email folder
- Target: Your clicks on the “THIS IS SPAM” button