

# ML4 – Optimisation

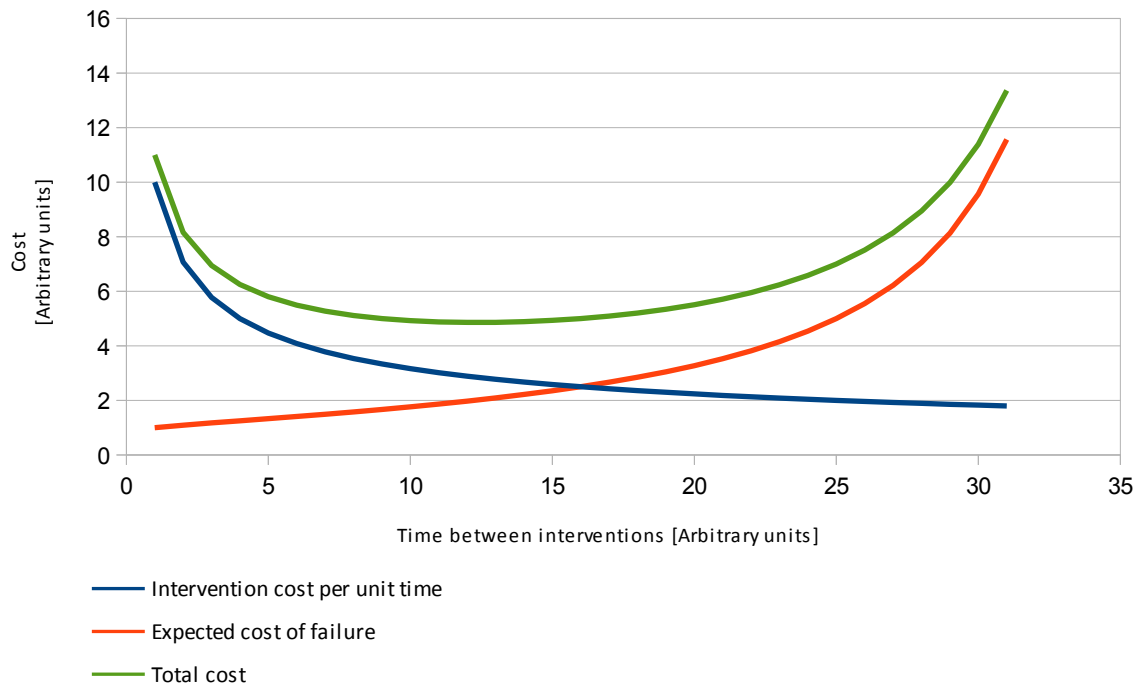
Stefano Rovetta

A.y. 2023-2024

# Motivation

- Maintenance interventions on a machine operating in a production line
- Scheduled maintenance interventions
- If done early, they reduce the probability of failure (but cost increases)
- If done late, they reduce the cost of interventions (but failures become more frequent)

## Cost of scheduled maintenance



- I have a rectangular sheet of wood from which I have to cut parts of non-rectangular shape. What is the layout that maximises the number of parts, or minimises the wasted material?
- Select the mix of goods (solid items) that maximises the value/cost ratio of a container to be shipped.
- Design a mechanical linkage that provides conversion of linear motion into a given 2-D trajectory, while minimising the number of links and joints.
- Decide the optimal number of years to keep your car before changing it.
- Given a set of different financial assets, each characterised by its expected return and its degree of uncertainty (variance), find the proportions of assets (portfolio mix) that maximises the total return.

# Optimisation concepts and terms

# The task of optimization:

Finding extrema of an **objective function**  $J(\mathbf{w})$ , where  $J : S \subset \mathbb{R}^m \rightarrow \mathbb{R}$ .

$S$  = feasible region, solution space, search space, feasible set

An **extremum** is a point  $\mathbf{w}^* \in S$  that may be a **maximum** or **minimum**.

A point  $\mathbf{w}^*$  is a minimum if there is a neighbourhood  $R \subseteq S$ , where the following holds:

$$J(\mathbf{w}) \geq J(\mathbf{w}^*) \quad \forall \mathbf{w} \in R$$

i.e.: A minimum is a point where  $J$  has a value smaller than in any other point in a given neighbourhood.

A point  $\mathbf{w}^*$  is a maximum if it is a minimum of  $-J$ , or if  $\leq$  is used.

Note: We will consider MINIMIZATION



- A minimum is **relative** if  $R \subset S$  strictly  
i.e., there is some other point in  $S$  (outside  $R$ ) where  $J$  has a smaller value than  $J(\mathbf{w}^*)$ .


A **relative minimum** is a minimum only in a neighbourhood (i.e. locally)

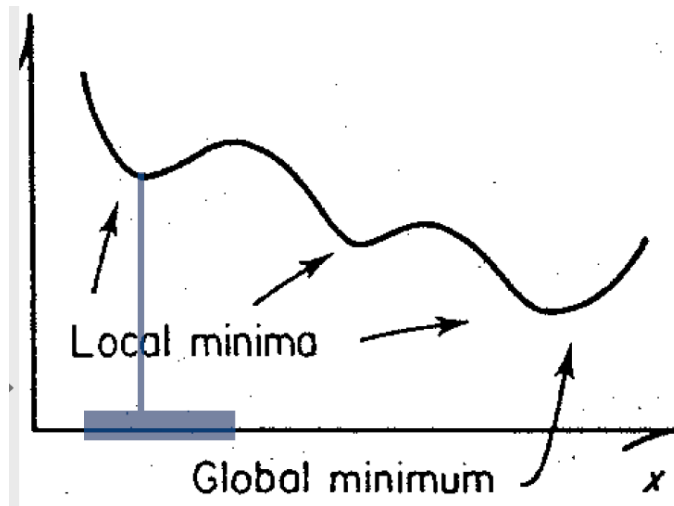
So we also use **local minimum**


- A minimum is **absolute** if  $R = S$ .


**Global minimum**

- A minimum is **isolated** if we can draw a sphere around  $\mathbf{w}^*$  with nonzero radius contained in  $R$   
(i.e. if there are no other minima adjacent to it)









# Objective functions in machine learning

- Often: **Cost function**, inversely related to how much we like a solution (the higher the cost, the less the quality of the solution)

- **Risk** (expected loss)

$$J(\mathbf{w}) = \int_{\mathcal{X}} \lambda(\mathbf{t}, \mathbf{y}(\mathbf{w})) p(\mathbf{x}) d\mathbf{x}$$

- **Regularised risk**

Expected loss + terms penalising “irregular” or “extreme” or “weird” solutions


Example: penalising high values of the parameters (**weight decay**)

$$J(\mathbf{w}) = \int_{\mathcal{X}} \lambda(\mathbf{t}, \mathbf{y}(\mathbf{w})) p(\mathbf{x}) d\mathbf{x} + \|\mathbf{w}\|$$

- **Likelihood** is a common objective that we have to **maximise** instead


- An **optimal solution** is an extremum  $\mathbf{w}^*$  of  $J$
- An **optimal value** is  $J(\mathbf{w}^*)$
- An **approximate solution with precision**  $\epsilon$  is a point  $\tilde{\mathbf{w}}$  whose value of the objective is close to an optimal solution:  $|J(\mathbf{w}^*) - J(\tilde{\mathbf{w}})| < \epsilon$   
absolute value not necessary once the type of problem (min or max) is given
- A **feasible solution** is any point  $\mathbf{w}$  which satisfies all hypotheses of the optimization problem

(a)




Minimum

(b)





Maximum

(c)



Saddle





## Convex sets

A set  $S \subset \mathbb{R}^m$  is convex if and only if, for any  $\theta \in [0, 1]$ ,

$$\forall \mathbf{v}, \mathbf{w} \in S \Rightarrow \theta \mathbf{v} + (1 - \theta) \mathbf{w} \in S$$

more generally if for any  $\theta_1 > 0, \dots, \theta_n > 0$  such that  $\sum_k \theta_k = 1$

$$\forall \mathbf{v}_1, \dots, \mathbf{v}_n \in S \Rightarrow \sum_k \theta_k \mathbf{v}_k \in S \quad \text{Convex combination}$$

Properties:

- $\mathbf{v} \in \mathbb{R}^m$  (a single point) is convex
- $\emptyset = \{ \}$  (the empty set) is convex
- $\mathbb{R}^m$  is convex

and if  $S_1$  and  $S_2$  are convex, then

- $S_1 \cap S_2$  is convex
- $S_1 \cup S_2$  IS NOT NECESSARILY convex ( $0 + 0 = \infty$ )

# Convex functions

A function  $J : S \subseteq \mathbb{R}^m \rightarrow \mathbb{R}$  is convex if  $S$  is a convex set and if  $\forall \mathbf{v}, \mathbf{w} \in S$ , and with  $0 \leq \theta \leq 1$ :


$$J(\theta \mathbf{v} + (1 - \theta) \mathbf{w}) \leq \theta J(\mathbf{v}) + (1 - \theta) J(\mathbf{w})$$


i.e., if its **epigraph** is a convex set

more generally if for any  $\theta_1 > 0, \dots, \theta_n > 0$  such that  $\sum_k \theta_k = 1$

$$\forall \mathbf{v}_1, \dots, \mathbf{v}_n \in S \Rightarrow J\left(\sum_k \theta_k \mathbf{v}_k\right) \leq \sum_k \theta_k J(\mathbf{v}_k)$$

$J$  is concave if  $-J$  is convex






## Convex functions, alternate (equivalent) definition

A function  $J : S \subseteq \mathbb{R}^m \rightarrow \mathbb{R}$  is convex if  $S$  is a convex set  
and if  $\forall \mathbf{v}, \mathbf{w} \in S$ ,

$$J(\mathbf{v}) \geq J(\mathbf{w}) + \nabla J(\mathbf{w}) \cdot (\mathbf{v} - \mathbf{w})$$



# Why should we be interested in convexity?

Convexity is a property of a problem, not just of a loss function:

- Objective
- Learning machine on which the objective is computed (because we want to take derivatives)
- **The parameter space** – remember that a function is not convex if its domain is not convex!

Convexity is a good thing:

- Uniqueness of extrema
- Convergence of iterative algorithms

Not all problems are convex:

- Some learners guarantee a convex problem (e.g., SVM)
- For non-convex problems we usually **cannot be sure whether a minimum is absolute (global) or relative (local)**

# Types of optimisation problems



## Types of optimisation problems – By type of optimisation variables $\mathbf{w}$

- **Discrete:**  $S \in \mathbb{N}^m$ , or more generally **countable variables**

Example: Maximise the number of boxes placed in a storage area

- **Continuous:**  $S \in \mathbb{R}^m$

Example: Minimise the surface area of a box or container to be built in expensive materials (e.g. rare wood)

## Types of optimisation problems – By feasible region $S$

- **Unconstrained:**  $S \equiv \mathbb{R}^m$  (or  $S \equiv \mathbb{N}^m$ )

Example: Find the coefficients of a linear-threshold classifier that minimise the average misclassification error ( $\approx$  error probability).

Feasible region  $S = \mathbb{R}^{d+1}$

- **Constrained:**  $S \in \mathbb{R}^m$  (or  $S \in \mathbb{N}^m$ )

Example: Given  $N$  financial assets, each with unit cost  $C_i$ , find the portfolio (quantity  $w_i$  per each asset) that maximises the expected return, *subject to*:

- $w_i \geq 0 \ \forall i : 1, \dots, N$  (non-negative quantities)
- $\sum_i C_i w_i \leq C_{\text{tot}}$  (the sum of invested money, sum of unit costs  $C_i$  times quantities  $w_i$ , must not exceed my total capital available!!)

Feasible region  $S$  a subset of the all-positive region  $\in \mathbb{R}^N$  (plus the origin)

Constraints are expressed by **constraint functions**:

- **Equality constraints:**  $f_1(\mathbf{w}) = 0, f_2(\mathbf{w}) = 0, \dots f_k(\mathbf{w}) = 0$
- **Inequality constraints:**  $f_1(\mathbf{w}) \geq 0, f_2(\mathbf{w}) \geq 0, \dots f_k(\mathbf{w}) \geq 0$

Both may be present in the same problem

## Types of optimisation problems – By smoothness of the objective

- **Smooth:** The objective is at least differentiable, maybe twice  
Example: Find the coefficient of a linear regression model in the least squares sense ( $\min J_{\text{mse}}$ )
- **Nonsmooth:** The objective is not differentiable  
Example: Find the coefficients of a linear-threshold classifier that minimise the average misclassification error ( $\approx$  error probability).

Obviously discrete  $\mathbf{w} \Rightarrow$  nonsmooth objective (it is not even continuous!)

## Types of optimisation problems – By shape of the objective

- **Linear**

Example: Buy the best car

- **Convex**

Example: Find the coefficient of a linear regression model in the least squares sense

- **Nonconvex**

Example: Find the pringing for goods in a shop, such that the expected income is maximised

- **special case: Quadratic**

Example: Find the coefficient of a linear regression model in the least squares sense

Easier to analyse but not guaranteed to be convex

Obviously linear  $\Rightarrow$  constrained, otherwise the objective can grow or decrease indefinitely!  
(Example: buy the best car with price  $\geq 0$  and price  $\leq$  my bank account balance)

## WARNING

The type of the loss alone **is not sufficient to describe the type of problem!**

$$\lambda(t, y) = \lambda(t, y(\mathbf{w}))$$

The loss as a function of the parameters is a **composite function**

loss  $\longrightarrow$  a function of  $t$  and  $y$

$y \longrightarrow$  a function of  $\mathbf{w}$

Examples:

Loss	Model		Problem
$\ t - y\ ^2$	Linear in $\mathbf{w}$ : $y(\mathbf{w}) = \mathbf{w} \cdot \mathbf{x}$	$\longrightarrow$	Quadratic
$\ t - y\ ^2$	Convex in $\mathbf{w}$ : $y(\mathbf{w}) = \log(1 + e^{\mathbf{w} \cdot \mathbf{x}})$	$\longrightarrow$	Convex, but not quadratic
$\ t - y\ ^2$	Non-convex in $\mathbf{w}$ : $y_i(\mathbf{w}) = e^{g_i(\mathbf{x}, \mathbf{w})} / \sum_j e^{g_j(\mathbf{x}, \mathbf{w})}$ , $i = 1 \dots c$	$\longrightarrow$	Non-convex

# The main types of problems in machine learning

- 1 Unconstrained, smooth, quadratic, convex problems

**Example:** least squares for linear models

- 2 Unconstrained, smooth, general (non-convex) problems

**Example:** all non-linear models, e.g. neural networks

- 3 Constrained, smooth, quadratic problems

**Example:** support vector machines, in general “kernel” methods

# Conceptual tools

## The gradient

$$\nabla J(\mathbf{w}) = \begin{bmatrix} \frac{\partial J(\mathbf{w})}{\partial w_1} \\ \frac{\partial J(\mathbf{w})}{\partial w_2} \\ \vdots \\ \frac{\partial J(\mathbf{w})}{\partial w_m} \end{bmatrix}$$


Other notation:

$\nabla_{\mathbf{w}} J(\mathbf{w})$       to highlight differentiation w.r.t.  $\mathbf{w}$

The gradient is a vector field (vector function of vector argument)



- Derivative  $\rightarrow$  rate of growth of a function of a scalar variable
- Negative sign  $\rightarrow$  decreasing
  
- Gradient *length (norm)*  $\rightarrow$  rate of maximum growth
- *Direction*  $\rightarrow$  *direction of maximum growth*



The gradient indicates the direction of maximum increase, and moving in the opposite direction  $-\nabla J(\mathbf{w})$  we achieve the *maximum rate of decrease*.

This observation is very useful in optimization techniques.

# Hessian matrix

or simply Hessian

$$H_J(\mathbf{w}) : \mathbb{R}^m \rightarrow \mathbb{R}^m \times \mathbb{R}^m \quad \text{s.t.} \quad H_J(\mathbf{w})_{ij} = \frac{\partial^2 J(\mathbf{w})}{\partial w_i \partial w_j}$$

Other notations:

$H$

$\nabla_{\mathbf{w}} \nabla_{\mathbf{w}} J(\mathbf{w})$       or       $\nabla_{\mathbf{w}}^2 J(\mathbf{w})$

$\nabla \nabla J(\mathbf{w})$       or       $\nabla^2 J(\mathbf{w})$

The Hessian matrix can be thought of as a list of  $m$  vectors

$$\mathbf{h}_i = \nabla \left( \frac{\partial J(\mathbf{w})}{\partial w_i} \right)$$

Derivative is a linear operator and the order of differentiation does not matter:

$$\frac{\partial^2 J(\mathbf{w})}{\partial w_i \partial w_j} = \frac{\partial}{\partial w_i} \left( \frac{\partial J(\mathbf{w})}{\partial w_j} \right) = \frac{\partial}{\partial w_j} \left( \frac{\partial J(\mathbf{w})}{\partial w_i} \right)$$

$\Rightarrow H$  is a symmetric matrix.

$$H = \begin{pmatrix} \frac{\partial^2 J}{\partial w_1^2} & \frac{\partial^2 J}{\partial w_1 \partial w_2} & \frac{\partial^2 J}{\partial w_1 \partial w_3} & \cdots & \frac{\partial^2 J}{\partial w_1 \partial w_m} \\ \frac{\partial^2 J}{\partial w_2 \partial w_1} & \frac{\partial^2 J}{\partial w_2^2} & \frac{\partial^2 J}{\partial w_2 \partial w_3} & \cdots & \frac{\partial^2 J}{\partial w_2 \partial w_m} \\ \frac{\partial^2 J}{\partial w_3 \partial w_1} & \frac{\partial^2 J}{\partial w_3 \partial w_2} & \frac{\partial^2 J}{\partial w_3^2} & \cdots & \frac{\partial^2 J}{\partial w_3 \partial w_m} \\ \vdots & & & & \\ \frac{\partial^2 J}{\partial w_m \partial w_1} & \frac{\partial^2 J}{\partial w_m \partial w_2} & \frac{\partial^2 J}{\partial w_m \partial w_3} & \cdots & \frac{\partial^2 J}{\partial w_m^2} \end{pmatrix}$$

# Taylor polynomials

The Taylor polynomial of degree 2 for a scalar function  $J(w)$  centered around  $w_0$ :

$$J(w) \approx J(w_0) + J'(w) \big|_{w=w_0} (w - w_0) + \frac{1}{2} J''(w) \big|_{w=w_0} (w - w_0)^2$$



Brook Taylor, 1685-1731

# Multi-dimensional Taylor polynomials

Equivalent formula when  $\mathbf{w} \in \mathbb{R}^m$ , centered around  $\mathbf{w}_0$ :

$$J(\mathbf{w}) \approx J(\mathbf{w}_0) + \nabla J(\mathbf{w})|_{\mathbf{w}=\mathbf{w}_0} (\mathbf{w} - \mathbf{w}_0) + \frac{1}{2}(\mathbf{w} - \mathbf{w}_0)^T H|_{\mathbf{w}=\mathbf{w}_0} (\mathbf{w} - \mathbf{w}_0)$$



# Characterizing minima

There are conditions for determining whether a feasible solution  $\mathbf{w}$  is a minimum


- One condition is true for all extrema (**necessary** condition of minimum, but not sufficient)
- One condition is true only for minima (**necessary and sufficient** condition of minimum)
- In the case of convex objectives, the necessary condition becomes sufficient

# Characterizing minima

Necessary first-order minimum condition:

$$\nabla J(\mathbf{w}^*) = 0$$

This condition characterizes all points which are local minima, but also local maxima or *saddle points* (points which are minima along one direction and maxima along another direction).




A saddle.



## The case of convex objective

The first-order condition is also **sufficient**!



An elliptic paraboloid (quadratic function with rotational symmetry) is a convex function


## Locally convex cost function

General, sufficiently smooth functions that are non-convex are nevertheless **convex in a neighbourhood of  $\mathbf{w}^*$**

The first-order condition is then a **necessary and sufficient condition of local minimum**. Other local minima belong to different neighborhoods (“basins”).

Need methods for searching in different basins:

- Multiple re-starts from different (random) initial points
- Occasional random “jumps” in addition to regular gradient descent
- Regularisation that makes basins appear gradually



An objective function that has two “basins”. The bottom of each basin is a local minimum.



# Definiteness and semidefiniteness

A matrix  $A$  is positive semidefinite (written  $A \succeq 0$ ) if  $\forall \mathbf{v} \in \mathbb{R}^m$

$$\mathbf{v}^T A \mathbf{v} \geq 0$$

It is positive definite (written  $A \succ 0$ ) if  $\forall \mathbf{v} \in \mathbb{R}^m$

$$\mathbf{v}^T A \mathbf{v} > 0$$

(similarly for “negative (semi)definite”)

Hessian generalizes the second derivative

“Positive semidefinite” generalises “non-negative”, “Positive definite” generalises “positive”

## Conditions of convexity

$J(\mathbf{w})$  is convex if  $H = \nabla^2 J(\mathbf{w})$  is **positive semidefinite** everywhere (for all  $\mathbf{w}$ )

$J(\mathbf{w})$  is locally convex around a point  $\mathbf{w}_0$  if  $H$  is positive semidefinite at  $\mathbf{w}_0$

To check whether a given point  $\mathbf{w}$  is a minimum...

- 1 Necessary conditions of extremum:  $\text{gradient} = 0$  in the point of interest
- 2 Necessary and sufficient condition of local convexity:  $\text{Hessian} \succeq 0$  at the point of interest
- 3 Sufficient conditions of local minimum:  $\text{gradient} = 0$  and  $\text{Hessian} \succeq 0$  at the point of interest
- 4 Necessary and sufficient condition of convexity:  $\text{Hessian} \succeq 0$  everywhere
- 5 Sufficient conditions of global minimum:  $\text{gradient} = 0$  and  $\text{Hessian} \succeq 0$  everywhere

# Basic optimisation methods

# The “oracle”

Important assumption:

In general, **we don't know everything about the objective function**  
(realistic problems are too complex!)

Example: The objective value is the result of a physical experiment

Example: The objective value is the performance of executing a program that is not open source

Example: The objective function depends on the data (training set) in a complex way

# The “oracle”

Given a point  $\mathbf{w}$  in the optimisation feasible space  $S$ ,  
we can query an **oracle**, a *black-box* system  
that can only provide information about some of these quantities:

- The value of the objective at  $\mathbf{w}$ ,  $J(\mathbf{w})$
- The gradient of the objective at  $\mathbf{w}$ ,  $\nabla J(\mathbf{w})$
- The Hessian of the objective at  $\mathbf{w}$ ,  $H = \nabla^2 J(\mathbf{w})$

# Types of optimization algorithms

- **Oracle of order 0:** we only know
  - $J(\mathbf{w})$       computed values of the objective
- **Oracle of order 1:** we know
  - $J(\mathbf{w})$
  - $\nabla J(\mathbf{w})$       computed values of the gradient
- **Oracle of order 2:** we know
  - $J(\mathbf{w})$
  - $\nabla J(\mathbf{w})$
  - and also  $H$       computed values of the Hessian

# A basic concept: Relaxation methods

## Relaxation sequence:

$$\mathbf{w}_0, \mathbf{w}_1, \mathbf{w}_2, \dots$$

such that

$$J(\mathbf{w}_0) \geq J(\mathbf{w}_1) \geq J(\mathbf{w}_2) \dots$$

## Approximation:

To generate a relaxation sequence, we employ **local approximations** to the objective, that are easier to deal with than the objective itself.

We will use **Taylor polynomials** of first and second order as our approximating functions.



A note about jargon:

- Where the mathematician see a sequence...
- ...the programmer sees a **loop** (iteration)

So “relaxation sequence”  $\equiv$  “iterative algorithm”

# The general relaxation algorithm

- Iteratively descend toward the minimum

$$\mathbf{w}_{i+1} = \mathbf{w}_i + \Delta \mathbf{w}_i$$

by taking *additive* steps  $\Delta \mathbf{w}_i$  chosen so as to decrease the objective  $J$


- Stop iterating when a pre-defined **stopping criterion** is satisfied

Examples of stopping criteria:

$$\nabla J(\mathbf{w}_i) < \epsilon \quad (\text{gradient} \approx 0)$$

$$J(\mathbf{w}_i) - J(\mathbf{w}_{i+1}) < \epsilon \quad (\text{improvement too small})$$

$$i == N_{\max} \quad (\text{max. number of iterations reached})$$



## Oracle of order 0: Direct search methods

We only know

- $J(\mathbf{w})$
- **NO TAYLOR EXPANSION AVAILABLE!**
- Only **direct search** techniques are possible:  
find promising points by using (meta)heuristics
  - Simulated annealing
  - genetic algorithms
  - particle swarm optimization
  - ant colony optimization
  - ...

or by branch-and-bound methods, or by random search...


PRO: simple, and guaranteed to find the GLOBAL extrema

CON: ...only if infinite time is available!

## Oracle of order 1: Gradient methods

The objective is locally approximated with its **first-order Taylor polynomial** around the current point

$$J(\mathbf{w}) \approx J(\mathbf{w}_i) + \nabla J(\mathbf{w}_i) \cdot (\mathbf{w} - \mathbf{w}_i)$$



# Gradient descent algorithm

- 1 Initialize: set  $i = 0$ ; select some  $\mathbf{w}_0 = \mathbf{w}_{\text{start}}$
- 2 Compute  $\nabla J(\mathbf{w}_i)$
- 3 Select the appropriate step size  $\eta_i$
- 4 Compute the step  $\Delta \mathbf{w}_i = -\eta_i \nabla J(\mathbf{w}_i)$
- 5 Perform step  $\mathbf{w}_{i+1} \leftarrow \mathbf{w}_i + \Delta \mathbf{w}_i$
- 6 Compute convergence test. If necessary, iterate from step 2.

```
epsilon = 1e-3;
maxiter = 100;
w = w0;
etavals = logspace(log10(.75), log10(.075), maxiter);
i = 1;
G = grad_J(w);
while norm(G) >= epsilon && i < maxiter
    eta = etavals(i);
    w = w - eta * G;
    G = grad_J(w);
end
```

## Step size strategies

- Constant  $\eta$ ,  $\eta_i = \eta_0 \quad \forall i$
- Predefined sequence of  $\eta_i$
- Full relaxation, or “line search”:  
$$\eta_i = \arg \min_{\eta} J(\mathbf{w}_i) + \eta \nabla J(\mathbf{w}_i) \quad \leftarrow \text{N.B. this is a one-dimensional function of } \eta$$
- Adaptive  $\eta$

# Features of gradient descent


## Pros


- Simple!


## Cons

- Unnecessarily slow convergence (always directed exactly as the negative gradient)











## Oracle of order 2: Second-order methods

The objective is locally approximated with its **second-order Taylor polynomial** around the current point


$$J(\mathbf{w}) \approx J(\mathbf{w}_i) + \nabla J(\mathbf{w}_i) \cdot (\mathbf{w} - \mathbf{w}_i) + (\mathbf{w} - \mathbf{w}_i)^T H(\mathbf{w}_i)(\mathbf{w} - \mathbf{w}_i)$$



# Newton-Raphson method



Isaac Newton (1642-1726)



Joseph Raphson (1648?-1715?)

Iterative relaxation method for finding **zeros** of functions.

# Newton-Raphson method

To find a **zero** of a function  $g(x)$ , iterate as follows:

- Start at  $x_0$
- At each step  $\tau$  compute the update as

$$x_{\tau+1} = x_{\tau} - \frac{g(x_{\tau})}{g'(x_{\tau})}$$

The update finds exactly the zero of the **1st order Taylor approximation of  $g$  in  $x_{\tau}$**   
...but the Taylor approximation is not  $g$ , so we repeat

# Newton-Raphson method for optimization

**Necessary 1st order condition of minimum:**

Minimum of  $J$  = zero of  $J'$

To find a **minimum** of a function  $J(x)$ , set  $g(x) = J'(x) = \frac{dJ(x)}{dx}$   
and then apply the Newton-Raphson method:

- Start at  $x_0$
- At each step  $\tau$  compute the update as

$$x_{\tau+1} = x_{\tau} - \frac{J'(x_{\tau})}{J''(x_{\tau})}$$

# Newton-Raphson method for multidimensional optimization

To find a **minimum** of a function  $J(\mathbf{w}) : \mathbb{R}^m \rightarrow \mathbb{R}$ ,

$$\mathbf{w}_{\tau+1} = \mathbf{w}_{\tau} - [H_J(\mathbf{w}_{\tau})]^{-1} \nabla J(\mathbf{w}_{\tau})$$

PRO: simple, much faster than gradient descent

CON: Need to compute the Hessian (  $O(m^2)$  space complexity)

and to invert it (  $O(n^q)$  time complexity, with  $2 < q \leq 3$  depending on the algorithm)



# Quasi-Newton methods

The most popular methods.

They use a **first-order oracle** (gradient only) to approximate **second-order information**.

Rationale:

Second derivative  $\approx (\text{first derivative at } \mathbf{w} + \mathbf{v} - \text{first derivative at } \mathbf{w}) / \|\mathbf{v}\|$

How do methods work in different problem types?

- Most problems in machine learning are **smooth**  $\Rightarrow$  **first- and second-order oracles**
- Many are unconstrained
- Many are non-convex

## Quadratic (convex) problems

- Encountered in linear regression
- First-order (gradient) and second-order (Newton, quasi-Newton) are efficient

# General (non-convex) problems

- Encountered in neural networks and in most non-trivial models
- First-order (gradient) and quasi-second-order can only **find local minima**
- Newton-type are usually impractical because  $H$  is too large
- To search for **better** minima, many optimisation cycles are attempted **starting from different initial conditions**
- This is a hybrid between random search (zero-order) and higher-order methods
- Several heuristics to speed up convergence exist:
  - adaptive step size (Vogl, SuperSAB, ... many others)
  - momentum
  - adaptive non-isotropic step size (AdaGrad, RMSprop, AdaDelta)
  - adaptive momentum (Adam, NAG-Nesterov's Accelerated Gradient, AdaMax, N-Adam)

# Example of acceleration heuristic

## Adaptive step size

The learning step size, or learning rate,  $\eta$  is changed according to some heuristics

Example:

```
if J increases:
    reduce eta
    if J increases more than a quantity T:
        cancel step and go back to previous iteration
else:
    increase eta
```

# Example of acceleration heuristic

## Momentum

The current step also includes a fraction  $\alpha < 1$  of the previous one

Example:

$$\Delta \mathbf{w}_i = -\eta \nabla J(\mathbf{w}_i) + \alpha \Delta \mathbf{w}_{i-1}$$

Helps traversing uninteresting regions where the gradient is small (“plateaus”, saddle points) by keeping a “memory” (inertia) of the previous motion.

# Constrained quadratic problems

- Encountered in support vector machines
- Methods for **constrained optimisation** are used to construct a dual problem, an unconstrained quadratic convex one
- Then methods for quadratic optimisation are used (first- and quasi-second order)